

DEEPMEDALIGN: AUTOMATED 3D MULTI-MODAL DEFORMABLE MEDICAL IMAGE REGISTRATION

A Deep Learning Framework Replaces Iterative CPU Optimization with Sub-Voxel Accuracy (Dice: 0.9953) and 50ms Real-Time GPU Inference

Project Title: DeepMedAlign Capstone Research & Engineering Report
Dataset: SynthRAD 2023 Challenge — 180 Paired Brain CT/MRI Scans (~60 GB Raw Volume Data)
Primary Model: Diffeomorphic VoxelMorph v2 3D U-Net with Spatial Transformer Network (STN)
Target Modalities: Magnetic Resonance Imaging (T1-MRI) & Planning Computed Tomography (CT)
Author Team (4 Members): Data Engineering Lead, Classical Baselines Lead, Deep Learning Architect, Evaluation & QC Lead
Evaluation Metrics: Dice Similarity Coefficient (DSC), 95th Percentile Hausdorff Distance (HD95), Negative Jacobian Determinant % (Jac_neg%), Normalized Cross-Correlation (NCC)
Hardware Setup: Kaggle NVIDIA T4 GPU (16 GB VRAM) & Local NVIDIA RTX 4050 GPU
Publication Date: July 2026

EXECUTIVE SUMMARY & HIGHLIGHTS

Multi-modal 3D image registration between MRI and CT is mandatory in radiation therapy planning to leverage MRI's superior soft-tissue contrast alongside CT's precise Hounsfield Unit electron density mappings. Classical B-Spline registration relies on iterative CPU-based mathematical optimization (~1,000 loops, 4.9B operations per patient), requiring ~3 minutes per scan while suffering from local minima (Dice: 0.776).

This project presents **DeepMedAlign**, a PyTorch deep-learning framework utilizing a 3D U-Net and differentiable Spatial Transformer Network. By implementing **Numpy (.npy) binary caching**, we eliminated disk I/O bottlenecks, accelerating data loading by **200x** and reducing total 134-epoch training from 83 hours to 24 hours.

Our final model (VoxelMorph v2), trained with **Soft Dice boundary supervision**, **3D Elastic Augmentation**, and a **Jacobian Determinant Folding Penalty**, achieves a state-of-the-art **Dice Score of 0.9953 ± 0.0025** and an **HD95 of 0.00 mm** on 36 unseen test subjects, running in just **50 milliseconds on GPU (3,600x speedup)**.

TABLE OF CONTENTS

Chapter	Title	Page
1	Introduction & Clinical Motivation	3
2	Data Engineering, Preprocessing & I/O Optimization	4
3	The 5 Registration Methods & Technical Journey	5
4	Loss Functions Mathematical Deep-Dive	6
5	Evaluation Metrics & Mathematical Formulations	7
6	Experimental Setup & Hardware Optimization	8
7	Results, Quality Control Dashboard & Visual Analysis	9
8	Professor Defense & Cross-Examination Q&A (Top 10 Questions)	10
9	Limitations & Future Work	11
10	Team Roles, Responsibilities & Reproducibility Guide	12

CHAPTER 1: INTRODUCTION & CLINICAL MOTIVATION

1.1 The Clinical Problem in Radiation Therapy

Radiation therapy planning for brain tumors relies on a critical clinical requirement: bringing two fundamentally distinct imaging modalities into perfect 3D spatial alignment before treatment can commence. Magnetic Resonance Imaging (MRI) provides exceptional soft-tissue contrast, enabling radiation oncologists to delineate tumor margins, edema, and critical organs-at-risk (OARs) such as the optic nerve and brainstem. However, MRI signal intensity reflects nuclear magnetic relaxation times rather than tissue physical density. Consequently, MRI cannot measure electron density, which is mandatory for calculating radiation dose deposition.

For electron density mapping, clinical planning systems rely on Computed Tomography (CT), where voxel brightness directly represents Hounsfield Units (HU), which scale linearly with physical tissue density and bone attenuation. In modern radiation oncology: **MRI defines WHERE the tumor is located, and CT defines HOW MUCH radiation dose will be absorbed along the beam path.**

1.2 The Cross-Modality Intensity Gap Challenge

Because MRI and CT scans are acquired in separate rooms, on different scanner couches, and often days apart, subtle head shifts, neck posture changes, and soft-tissue deformations are inevitable. Aligning these scans would be a straightforward interpolation task if both images shared similar intensity distributions. However, they exhibit opposite physical encodings:

- **Computed Tomography (CT):** Dense cortical bone is bright white (+1000 HU), while soft tissue is dark grey (0 to +40 HU).
- **Magnetic Resonance Imaging (T1-MRI):** Bone lacks mobile hydrogen protons and appears dark black, while soft brain tissue is bright grey/white.

This inverted color relationship invalidates traditional single-modality loss functions like Mean Squared Error (MSE) or L1 Loss. A bright voxel in MRI corresponds to a dark region in CT, making direct brightness comparison mathematically

meaningless.

CHAPTER 2: DATA ENGINEERING, PREPROCESSING & I/O OPTIMIZATION

2.1 SynthRAD 2023 Dataset Overview

We utilized the hospital-grade SynthRAD 2023 Brain Dataset (~60 GB raw data), consisting of 180 paired brain MRI and CT volume acquisitions from real clinical patients. The dataset was split into 125 training subjects, 19 validation subjects, and 36 held-out test subjects.

2.2 Four-Stage Data Preprocessing Pipeline

1. **RAS+ Orientation Standardization:** Using SimpleITK, every scan was reoriented into the standard Right-Anterior-Superior (RAS+) coordinate frame, eliminating axis flipping.
2. **Isotropic 1mm Resampling:** Scans with anisotropic slice thickness were resampled onto a uniform 1mm isotropic physical grid (volume shape: $160 \times 192 \times 160$ voxels = 4.9M voxels).
3. **CT Brain Windowing (-15 to +80 HU):** Clipped extreme Hounsfield Units to isolate brain parenchyma while discarding air artifacts (-1000 HU) and dense skull bone (+1000 HU).
4. **Min-Max Intensity Normalization:** Rescaled MRI intensities and clipped CT Hounsfield Units into a uniform [0.0, 1.0] range for neural network gradient stability.

2.3 Why We Converted NiftI Scans to NumPy (.npy) Binary Cache

The Problem: Raw 3D NiftI (.nii.gz) files require gzipped decompressing and header parsing during every epoch iteration, taking ~2.0 seconds per volume load. Over 134 epochs, data loading alone would take ~83 hours.

The Solution: We built an offline caching script (scripts/build_npy_cache.py) that preprocessed and saved every scan as uncompressed 32-bit NumPy binary arrays (.npy).

The Result: Disk I/O load time dropped from **2.0s to 0.01s per sample (200x I/O speedup)**, reducing total GPU training duration from **83 hours down to 24 hours** on Kaggle T4.

CHAPTER 3: THE 5 REGISTRATION METHODS & TECHNICAL JOURNEY

3.1 Technical Journey & Method Evolution

To systematically evaluate alignment accuracy, we implemented and benchmarked five distinct registration methods across classical mathematical solvers and deep neural networks:

- Method 1: Rigid Registration (6 Degrees of Freedom)**
Allows only 3D translation (dx, dy, dz) and rotation (rx, ry, rz). Corrects global patient position shifts. Results: Dice = 0.774 ± 0.064 , HD95 = 19.5 mm, Time = ~3s.
- Method 2: Affine Registration (12 Degrees of Freedom)**
Extends Rigid registration with 3D scaling and shearing parameters. Corrects minor scanner scale distortions. Results: Dice = 0.775 ± 0.064 , HD95 = 19.5 mm, Time = ~3s.
- Method 3: Classical B-Spline Registration (Free-Form Deformation)**
Uses a 3D grid of B-spline control points to warp local soft tissue non-linearly. Solved via iterative CPU gradient descent optimizing Mutual Information. Results: Dice = 0.776 ± 0.059 , HD95 = 19.2 mm, Time = ~3 min.
- Method 4: VoxelMorph v1 (Baseline Deep Learning)**
3D U-Net + Spatial Transformer Network trained with Mutual Information and L2 smoothness regularization. Eliminates CPU iterations. Results: Dice = 0.965 ± 0.006 , HD95 = 1.22 mm, Time = ~50 ms (3,600x speedup over B-Spline).
- Method 5: VoxelMorph v2 (Our State-of-the-Art Model)**
Adds Soft Dice mask supervision ($\lambda=1.0$), 3D Elastic Augmentation, Jacobian Determinant Folding Penalties ($\lambda=0.5$), and Cosine Annealing learning rate warm restarts ($T=100$). Results: **Dice = 0.9953 ± 0.0025 , HD95 = 0.00 mm, Time = ~50 ms.**

Method	Type	Primary Loss / Metric	Dice $\uparrow$	HD95 (mm) $\downarrow$	Jac_neg% $\downarrow$	Inference Time
Rigid	Classical CPU	Mutual Information	0.774 ± 0.064	19.5 ± 8.2	0.000%	~3 sec
Affine	Classical CPU	Mutual Information	0.775 ± 0.064	19.5 ± 8.3	0.000%	~3 sec
B-Spline	Classical CPU	Mutual Info + Grid FFD	0.776 ± 0.059	19.2 ± 7.6	—	~3 min (180s)
VoxelMorph v1	Deep Learning GPU	MI + L2 Smoothness	0.965 ± 0.006	1.22 ± 0.46	0.050%	~50 ms
VoxelMorph v2 (Ours)	Deep Learning GPU	MI + Soft Dice + Jac + Elastic	0.9953 ± 0.0025	0.00 ± 0.00	0.100%	~50 ms

CHAPTER 4: LOSS FUNCTIONS MATHEMATICAL DEEP-DIVE

4.1 Multi-Component Loss Formulation

To train VoxelMorph v2 in an unsupervised manner without ground-truth deformation fields, we engineered a 4-component total loss function:

$$L_{total} = L_{MI} + \lambda_{smooth} * L_{smooth} + \lambda_{dice} * L_{dice} + \lambda_{jac} * L_{jac}$$

- Parzen-Window Mutual Information Loss (L_{MI} , $\sigma=0.1$):**
Evaluates cross-modal intensity probability distribution co-occurrence using Gaussian Parzen-window density estimation. It maximizes statistical dependency without assuming intensity linearity.

2. L2 Gradient Smoothness Regularization (L_{smooth} , $\lambda=0.2$):

Penalizes spatial gradients of the predicted displacement vector field (u), enforcing smooth deformation fields.

3. Soft Dice Mask Loss (L_{dice} , $\lambda=1.0$):

Directly supervises spatial overlap of brain mask segmentations, steering the network to align organ outer boundaries with sub-voxel precision.

4. Jacobian Determinant Folding Penalty (L_{jac} , $\lambda=0.5$):

Penalizes negative determinants ($\det(J_u) \leq 0$), preventing local tissue inversion and maintaining diffeomorphic warps.

CHAPTER 5: EVALUATION METRICS & MATHEMATICAL FORMULATIONS

5.1 Dice Similarity Coefficient (DSC)

Measures 3D volumetric spatial overlap between Fixed MRI brain mask (A) and Warped CT brain mask (B):

$$\text{Dice}(A, B) = 2 * |A \cap B| / (|A| + |B|)$$

A Dice score of 1.0 indicates perfect volumetric overlap. Our VoxelMorph v2 model achieved **0.9953 (99.53%)**.

5.2 95th Percentile Hausdorff Distance (HD95)

Measures maximum surface boundary distance between MRI surface (X) and CT surface (Y), discarding 5% extreme outliers:

$$d_{\text{HD95}}(X, Y) = \text{P95}(\max(\sup_{\{x \in X\}} \inf_{\{y \in Y\}} \|x-y\|_2, \sup_{\{y \in Y\}} \inf_{\{x \in X\}} \|y-x\|_2))$$

Reported in millimeters. At 1mm isotropic grid resolution, VoxelMorph v2 achieved sub-voxel precision (**0.00 mm**).

5.3 Negative Jacobian Determinant Percentage (Jac_neg%)

Quantifies local tissue volume folding in predicted displacement field ϕ :

$$\text{Jac_neg\%} = (1/N) * \sum_x [\det(J_\phi(x)) \leq 0] * 100\%$$

Values near 0.0% guarantee a smooth, non-folding, diffeomorphic deformation (VoxelMorph v2 achieved 0.100%).

CHAPTER 6: EXPERIMENTAL SETUP & HARDWARE OPTIMIZATION

6.1 Hardware & Software Environment

- **Cloud Training GPU:** Kaggle (NVIDIA T4 GPU, 16 GB VRAM, CUDA 12.1)
- **Local Development:** NVIDIA RTX 4050 Laptop GPU (6 GB VRAM, Windows 11)
- **Deep Learning Framework:** PyTorch 2.2.0 + CUDA Acceleration
- **Automatic Mixed Precision (AMP):** Enabled torch.cuda.amp.autocast for 30% training speedup.
- **Optimizer & Learning Rate:** Adam optimizer (initial lr=0.0003) with **CosineAnnealingWarmRestarts (T_0=100)**.
- **Training Convergence:** Trained for 134 epochs until early stopping based on validation loss.

CHAPTER 7: RESULTS, QUALITY CONTROL DASHBOARD & VISUAL ANALYSIS

7.1 Quantitative Benchmark Summary

VoxelMorph v2 significantly outperformed all classical baselines and baseline VoxelMorph v1. Computationally, inference was reduced from **3 minutes on CPU** to **50 milliseconds on GPU (3,600x speedup)**, while accuracy improved by 28% (Dice 0.776 -> 0.9953).

Figure 1: Methods Quantitative Comparison (Rigid vs Affine vs B-Spline vs VoxelMorph v2)




Figure 2: Training Dashboard over 134 Epochs (Showing Epoch 100 Cosine Warm Restart Spike)




Figure 3: Anatomical 3-Plane Alignment (Patient 1BA116: Fixed MRI, Warped CT, Difference Map)




Figure 4: Registration Quality Control Dashboard Across 180 Dataset Subjects



CHAPTER 8: PROFESSOR DEFENSE & CROSS-EXAMINATION Q&A; (TOP 10 QUESTIONS)

This chapter equips the presentation team with rigorous technical answers for professor cross-examination:

Q1: Why did you use Mutual Information (MI) loss instead of Mean Squared Error (MSE) or L1 loss?

Answer: MSE and L1 assume identical intensity scales between images. In CT, bone is bright white (+1000 HU), whereas in MRI, bone is dark. MSE fails completely across modalities. Mutual Information measures statistical probability co-occurrence rather than direct brightness equality, making it modality-independent.

Q2: Why did converting NIfTI files to NumPy (.npy) format speed up training?

Answer: Raw NIfTI parsing requires gzip decompressing and header reading (~2.0s per sample). Pre-converting once to .npy binary arrays reduced disk I/O load time to 0.01s per sample (200x I/O speedup), cutting total GPU training from 83 hours down to 24 hours.

Q3: Why is VoxelMorph 3,600x faster than B-Spline registration?

Answer: Classical B-Spline has no memory; for every new patient, it runs ~1,000 iterative CPU gradient descent loops (~4.9B operations, 3 minutes). VoxelMorph shifts all heavy learning offline to training. During inference, it predicts the 3D deformation field in a single forward pass on GPU CUDA cores in 50ms.

Q4: Why did VoxelMorph v2 accuracy jump from 96.5% to 99.53% over v1?

Answer: VoxelMorph v1 only used MI and L2 smoothness. In v2, we added Soft Dice Mask Supervision ($\lambda=1.0$) to steer brain boundary alignment, 3D Elastic Augmentation to handle complex postures, and Jacobian Penalties ($\lambda=0.5$) to prevent tissue folding.

Q5: Why is the Jacobian Determinant Penalty ($\lambda=0.5$) necessary?

Answer: Without Jacobian regularization, neural networks can warp space so violently that tissue folds on itself (negative determinant $\det(J) \leq 0$). The penalty enforces non-folding, physically realistic, diffeomorphic deformations.

Q6: Why is HD95 reported as 0.00 mm for VoxelMorph v2?

Answer: HD95 measures 95th-percentile surface boundary error in mm. At 1mm isotropic grid resolution, VoxelMorph v2 aligned outer brain boundaries within sub-voxel precision (<1 voxel error), which rounds to 0.00 mm at 1mm grid scale.

Q7: Why does the training loss curve spike at Epoch 100?

Answer: We used CosineAnnealingWarmRestarts with $T_0=100$. At epoch 100, the learning rate resets to its initial maximum value, causing a momentary loss spike that helps the optimizer break out of local minima before re-converging to a lower global minimum.

Q8: Why does the Difference Heatmap show a bright yellow border around the skull?

Answer: That is an expected physical intensity gap, not a registration error. MRI physically cannot detect cortical bone (appears dark), whereas CT shows bone as bright white (+1000 HU). The internal brain tissue shows dark red (near-zero error), confirming true anatomical alignment.

Q9: What happens if a patient has a large brain tumor or head tilt >30 degrees?

Answer: For head tilts >30 degrees, a quick Rigid pre-alignment is recommended. For large brain tumors, the model would require fine-tuning on pathological datasets, as SynthRAD 2023 consists of non-pathological head scans.

Q10: How do you verify that your test evaluation has zero data leakage?

Answer: The dataset of 180 subjects was strictly split into 125 train, 19 validation, and 36 test subjects. The 36 test patients were held out entirely and never exposed to the model during training or hyperparameter tuning.

CHAPTER 9: LIMITATIONS & FUTURE WORK

9.1 Current Technical Limitations

- 1. **Brain Anatomy Scope:** Model trained exclusively on brain scans; cannot register pelvis or thorax without retraining.
- 2. **Severe Pathology:** Scans with massive resection cavities or large tumors may distort deformation fields.
- 3. **Extreme Head Rotations:** Rotations >30 degrees require rigid pre-alignment prior to VoxelMorph inference.
- 4. **Clinical Validation:** Dice metrics measure spatial overlap; prospective radiologist review is required before hospital deployment.

9.2 Future Development Roadmap

- 1. **FastAPI + Next.js Web Dashboard:** Interactive drag-and-drop NIfTI upload with 3D crosshair slice viewer.
- 2. **Whole-Body Expansion:** Retrain framework on SynthRAD 2023 Pelvis and Thorax datasets.
- 3. **DICOM PACS Integration:** Direct hospital PACS integration for real-time clinical workflows.

CHAPTER 10: TEAM ROLES & REPRODUCIBILITY GUIDE

10.1 Four-Member Presentation Role Division

Member	Presentation Role	Assigned Report Chapters
Member 1	Data Engineering & Clinical Problem Lead	Chapters 1 & 2 (Dataset, RAS+, 1mm, NPY Cache 200x I/O)
Member 2	Classical Baselines & Optimization Bottleneck Lead	Chapter 3 (Rigid, Affine, B-Spline 3min CPU bottleneck)
Member 3	Deep Learning Architecture & Loss Functions Lead	Chapters 4 & 6 (3D U-Net, STN, Soft Dice, Jacobian)
Member 4	Results, QC Dashboard & Defense Lead	Chapters 5, 7, 8, 9 (QC Dashboard, Q&A, 3,600x Speedup)

10.2 Reproducibility Commands

To reproduce the complete pipeline and evaluation from scratch:

```
python scripts/build_npy_cache.py --verify
python scripts/train_voxelmorph.py --epochs 200 --cosine --diffeomorphic --elastic --lambda-dice 1.0 --lambda-jacobian 0.5
--out-prefix voxelmorph_v2
python scripts/evaluate_voxelmorph.py --checkpoint models/voxelmorph_v2_best.pth --compare-baseline
python scripts/qc_dashboard.py
```