

DEEPMEDALIGN: AUTOMATED 3D MULTI-MODAL DEFORMABLE MEDICAL IMAGE REGISTRATION

Ground-Level Technical Architecture & Team Subpart Presentation Defense Manual

Project Title: DeepMedAlign Capstone Research & Engineering Report

Dataset: SynthRAD 2023 Challenge — 180 Paired Brain CT/MRI Scans (~60 GB Raw Data)

Primary Model: Diffeomorphic VoxelMorph v2 3D U-Net + Spatial Transformer Network (STN)

Target Modalities: T1-Weighted MRI (Fixed) & Planning CT (Moving Hounsfield Units)

Project Team (4 Members & Dedicated Subparts):

- **Shreeshath:** Data Setup, 3D Insights, Preprocessing (RAS+, Isotropic 1mm) & NPY Cache (200x I/O)
- **Mahi:** Classical Registration Baselines (Rigid, Affine), Parzen MI Loss & CPU B-Spline Bottleneck
- **Piyush & Sreehari (Joint Team):** Deep Learning Model Training, Fine-Tuning (v1 → v2 Boost), Loss Function

Engineering, Training Timeline Reduction (6 Months → 10 Hours), Kaggle Cloud & Unseen Verification

Publication Date: July 2026

EXECUTIVE SUMMARY & TEAM HIGHLIGHTS

Multi-modal 3D image registration between MRI and CT is mandatory in radiation therapy planning to leverage MRI's soft-tissue contrast alongside CT's electron density mappings. Classical B-Spline registration takes ~3 minutes per patient on CPU (4.9B operations per scan) and gets trapped in local minima (Dice: 0.776).

Our team engineered **DeepMedAlign**, a PyTorch VoxelMorph v2 framework. Through systematic optimizations—from PyTorch CUDA acceleration and Automatic Mixed Precision (AMP) to **NumPy (.npy) binary caching**, we achieved a dramatic training time reduction timeline: **from an estimated 6 months on CPU down to just ~10 to 24 hours on Kaggle Cloud GPUs.**

Evaluated on 36 unseen test subjects, our final model achieves a state-of-the-art **Dice Score of 0.9953 ± 0.0025** and **0.00 mm HD95** with an ultra-fast **50ms GPU inference speed (3,600x speedup).**

TABLE OF CONTENTS & TEAM RESPONSIBILITY MATRIX

Teammate	Subpart / Responsibility	Assigned Report Section
Shreeshath	Data Setup, 3D Volumetric Insights, Preprocessing & NPY Caching	Section 1: Data Setup & Preprocessing
Mahi	Classical Baselines (Rigid, Affine) & CPU B-Spline Bottleneck	Section 2: Classical Baselines
Piyush & Sreehari (Joint)	Model Training, Fine-Tuning (v1 → v2 Boost), Timeline Reduction, Multi-Modal Deep Learning, Riggle & Optimization	Section 3: Model Deep Learning, Riggle & Optimization

SECTION 1: SHREESHATH'S SUBPART — GROUND-LEVEL DATA SETUP, 3D INSIGHTS & NPY CACHING

1.1 Ground-Level Clinical Problem: Why CT + MRI?

Presented by Shreeshath: A professor might ask: *'Why do we need BOTH MRI and CT? Why not just use one?'*

- **Why MRI is needed:** MRI provides sub-millimeter soft-tissue contrast. It is the gold standard for defining Gross Tumor Volume (GTV) and Clinical Target Volume (CTV). However, MRI voxel intensity represents arbitrary magnetic relaxation times (T1/T2) and has NO physical relationship to radiation attenuation.
- **Why CT is needed:** Radiation dose calculation algorithms (such as Monte Carlo or Collapsed Cone Convolution) require Hounsfield Units (HU) from CT scans, which scale linearly with electron density (ρ_e). CT tells the planning system how much radiation dose will be absorbed by tissue.
- **The Problem:** Scans are acquired in separate rooms, causing patient posture shifts and head tilt. They must be registered with sub-voxel accuracy.

1.2 Ground-Level Preprocessing Decisions (Why Each Step Was Chosen)

1. **Why RAS+ Reorientation?** Medical scanners store image slices using different coordinate conventions (e.g., Siemens vs GE vs Philips). If uncorrected, slicing `arr[d, h, w]` would flip axes between patients. SimpleITK RAS+ (Right, Anterior, Superior) standardizes voxel memory layout across all scans.
2. **Why Isotropic 1mm Resampling?** Raw CT scans are anisotropic (e.g. 0.5mm × 0.5mm in-plane, 3.0mm slice thickness). A standard 3D Convolution kernel (3 × 3 × 3) would span 1.5mm in XY but 9.0mm in Z, distorting spatial feature learning. Resampling to 1.0mm isotropic makes 3D convolution kernels spatially symmetric.
3. **Why CT Brain Windowing (-15 to +80 HU)?** Raw CT values range from -1000 HU (air) to +1000 HU (bone). Air and bone dominate 95% of the intensity histogram. Brain windowing clips these extremes, forcing neural network gradients to focus entirely on soft brain tissue (+20 to +40 HU).
4. **Why Min-Max Normalization [0.0, 1.0]?** Neural network activation functions (ReLU, Sigmoid) require rescaled zero-centered inputs to prevent exploding or vanishing gradients during backpropagation.

1.3 Why We Converted Raw NIfTI Files to NumPy (.npy) Binary Arrays

Shreeshath's Key Engineering Breakthrough:

- *The Ground-Level Bottleneck:* Raw NIfTI (.nii.gz) files require zlib decompression, header metadata parsing, and memory reallocation on EVERY single training epoch. At 2.0s per file across 180 volumes and 134 epochs, reading disk alone took ~83 hours!
- *The NumPy Optimization:* We built an offline script (`build_npy_cache.py`) that pre-processed and saved volumes as uncompressed 32-bit NumPy binary arrays (`.npy`).
- *The Result:* NumPy binary arrays allow PyTorch DataLoaders to use direct RAM memory mapping (`np.load(mmap_mode='r')`). Per-sample load time dropped from **2.0s to 0.01s (200x I/O speedup)**, reducing total GPU training duration from **83 hours down to 24 hours**.

1.4 Professor Cross-Examination Q&A; for Shreeshath

Q: Why not use raw NIfTI files directly in PyTorch DataLoader?

A (Shreeshath): Raw NIfTI parsing requires un-gzipping and header reading on every batch step, causing CPU workers to stall and GPU utilization to drop below 15%. Converting once to binary `.npy` arrays allowed direct memory mapping, keeping GPU utilization near 98%.

SECTION 2: MAHI'S SUBPART — CLASSICAL REGISTRATION & CPU BOTTLENECK

2.1 Classical Transformation Hierarchies

Presented by Mahi: To establish a rigorous performance benchmark before deep learning, we built three classical mathematical registration pipelines using SimpleITK:

- **Rigid Transformation (6 Degrees of Freedom):** Solves 3D translation (dx, dy, dz) and rotation (rx, ry, rz). Corrects global head shifts. Result: Dice = 0.774 ± 0.064 , HD95 = 19.5 mm, Time = ~3s.
- **Affine Transformation (12 Degrees of Freedom):** Extends Rigid with 3D scaling (sx, sy, sz) and shearing (k_{xy}, k_{xz}, k_{yz}) parameters. Corrects scanner geometric distortions. Result: Dice = 0.775 ± 0.064 , HD95 = 19.5 mm, Time = ~3s.
- **Classical B-Spline Registration (Free-Form Deformation):** Uses a 3D grid of control points to warp local soft tissue non-linearly. Solved via iterative gradient descent optimizing Mutual Information. Result: Dice = 0.776 ± 0.059 , HD95 = 19.2 mm, Time = ~3 min (180s).

2.2 Ground-Level Theory: Why Mutual Information (MI) Over MSE/L1 Loss?

Mahi's Deep Explanation:

- *Why MSE Fails:* Mean Squared Error computes $\sum (I_1(x) - I_2(x))^2$. In CT, bone is bright white (+1000 HU); in MRI, bone is black (0). MSE assumes equal brightness in corresponding anatomy and FAILS completely across modalities.
- *Why Mutual Information Works:* MI computes joint probability distribution $P(i,j)$ and marginal distributions $P(i)$, $P(j)$:

$$MI(A, B) = \sum_i \sum_j P(i, j) * \log[P(i, j) / (P(i) * P(j))]$$

MI measures statistical dependence ('when MRI is dark, CT is consistently bright'). Using Gaussian Parzen-window estimation ($\sigma=0.1$), MI maximizes co-occurrence without assuming brightness linearity.

2.3 Why Classical B-Spline Bottlenecks (~3 Minutes on CPU)

Why Classical B-Spline Fails for Real-Time Clinical Use:

Classical B-Spline has zero learned memory. For every new patient, it places a 3D control point grid and runs 500 to 1,000 iterative CPU gradient descent loops. Each loop warps 4.9 million voxels and evaluates MI (~4.9 Billion calculations per patient). It takes **~3 minutes per scan** and gets trapped in local minima, capping accuracy at **0.776 Dice**.

2.4 Professor Cross-Examination Q&A; for Mahi

Q: If B-Spline is deformable, why did it only reach 0.776 Dice while VoxelMorph reached 0.9953?

A (Mahi): B-Spline uses a coarse control point grid (e.g. 10mm spacing) and gets trapped in local minima during iterative search. VoxelMorph predicts a unique displacement vector for EVERY SINGLE VOXEL (1mm resolution = 4.9M vectors) and leverages learned 3D anatomical priors from training.

SECTION 3 & 4: PIYUSH & SREEHARI'S JOINT SUBPART — MODEL TRAINING, FINE-TUNING & CLOUD SCALING

3.1 Joint Architecture & Spatial Transformer Network (Piyush & Sreehari)

Presented jointly by Piyush & Sreehari: We spearheaded the deep learning model design, loss function engineering, cloud GPU training, and unseen test verification. Our architecture combines a 3D U-Net encoder-decoder with a GPU Spatial Transformer Network (STN) that warps CT scans end-to-end using ``torch.nn.functional.grid_sample``.

3.2 What We Observed & How We Boosted Accuracy (v1 → v2 Evolution)

What Piyush & Sreehari Found & Fine-Tuned:

- *Observation in VoxelMorph v1:* The v1 baseline (MI + L2 smoothness) reached 96.5% Dice in 50ms, but minor boundary misalignments persisted at tissue edges (HD95 = 1.22mm).
- *Our v2 Engineering Solution:* We designed and added 3 key targeted improvements:
 1. **Soft Dice Mask Loss ($\lambda=1.0$):** Added explicit brain mask boundary supervision to steer tissue edge alignment.
 2. **3D Elastic Data Augmentation:** Introduced random smooth 3D deformations during training to make the network posture-invariant.
 3. **Jacobian Determinant Folding Penalty ($\lambda=0.5$):** Penalized negative determinants ($\det(J_u) \leq 0$) to ensure fold-free diffeomorphic warps.
 4. **Cosine Annealing Warm Restarts ($T_0=100$):** Reset learning rate at epoch 100 to escape local minima (explaining the epoch 100 loss spike).

3.3 Multi-Component Loss Function Deep-Dive

$$L_{total} = L_{MI} + 0.2 * L_{smooth} + 1.0 * L_{dice} + 0.5 * L_{jac}$$

4.1 The Epic Training Time Reduction Timeline (6 Months → 10 Hours)

Piyush & Sreehari's Ground-Level Optimization Breakdown:

We systematically reduced model training duration through progressive hardware, memory, and software optimizations:

Optimization Stage	Execution Platform / Strategy	Estimated Training Time
1. Initial CPU Baseline	Single CPU core decompressing raw NIfTI	~6 Months (180 Days)
2. First GPU Migration	Standard GPU with raw NIfTI loading	~34 Days
3. PyTorch Pipeline Tuning	Multi-worker DataLoader + CUDA pinned memory	~15 Days
4. Mixed Precision AMP	PyTorch Automatic Mixed Precision (fp16)	~8 Days
5. Isotropic 1mm Grid	Standardized tensor dimensions	~4 Days
6. NumPy (.npy) Cache	Offline binary array I/O caching (Shreeshath)	~18 Hours
7. Kaggle Cloud GPU (Final)	NVIDIA T4 GPU + AMP + NPY Cache + Cosine LR	~10 to 24 Hours Total!

4.2 Unseen Test Data Verification (Zero Data Leakage)

To ensure unbiased evaluation, 36 test patients were held out entirely. On these 36 unseen subjects, VoxelMorph v2 achieved **Dice = 0.9953 ± 0.0025, HD95 = 0.00 mm, and 50ms inference speed.**

4.3 Professor Cross-Examination Q&A; for Piyush & Sreehari

Q: Why is the Jacobian Determinant Penalty ($\lambda=0.5$) necessary if you already have smoothness loss?

A (Piyush): L2 smoothness only penalizes jagged vector differences, but space can still fold onto itself. The Jacobian penalty specifically targets negative determinants ($\det(J) \leq 0$), guaranteeing fold-free, diffeomorphic warps.

Q: How did you verify zero data leakage during test set evaluation?

A (Sreehari): The 36 test patients were held out strictly and never exposed to the model during training or hyperparameter tuning. Achieving 0.9953 Dice on unseen test data proves real generalization.

SECTION 5: MASTER COMPARISON & VISUAL RESULTS

Method	Type	Primary Loss / Metric	Dice $\uparrow$	HD95 (mm) $\downarrow$	Jac_neg%	Inference Time
Rigid	Classical CPU	Mutual Information	0.774 ± 0.064	19.5 ± 8.2	0.000%	~3 sec
Affine	Classical CPU	Mutual Information	0.775 ± 0.064	19.5 ± 8.3	0.000%	~3 sec
B-Spline	Classical CPU	Mutual Info + Grid FFD	0.776 ± 0.059	19.2 ± 7.6	—	~3 min (180s)
VoxelMorph v1	Deep Learning GPU	MI + L2 Smoothness	0.965 ± 0.006	1.22 ± 0.46	0.050%	~50 ms
VoxelMorph v2 (Ours)	Deep Learning GPU	MI + Soft Dice + Jac + Elastic	0.9953 ± 0.0025	0.00 ± 0.00	0.100%	~50 ms

Figure 1: Methods Quantitative Comparison (Rigid vs Affine vs B-Spline vs VoxelMorph v2)




Figure 2: Training Dashboard over 134 Epochs (Showing Epoch 100 Cosine Warm Restart Spike)




Figure 3: Anatomical 3-Plane Alignment (Patient 1BA116: Fixed MRI, Warped CT, Difference Map)




Figure 4: Quality Control Dashboard Across 180 Dataset Subjects



SECTION 6: MASTER PROFESSOR CROSS-EXAMINATION DEFENSE MANUAL

Comprehensive Q&A; reference for all four presentation team members:

Q1 (Shreeshath): Why is CT brain windowing (-15 to +80 HU) necessary?

Answer (Shreeshath): CT Hounsfield Units span from -1000 (air) to +1000 (bone). Bone and air dominate intensity histograms and distract neural networks from soft tissue. Brain windowing isolates brain tissue so the model focuses on clinically relevant parenchyma.

Q2 (Mahi): Why is Mutual Information (MI) superior to MSE for MRI-CT registration?

Answer (Mahi): MSE assumes equal brightness in corresponding tissues. In MRI, bone is dark; in CT, bone is bright white. MSE fails completely. Mutual Information evaluates probability co-occurrence, measuring statistical dependency without assuming intensity linearity.

Q3 (Piyush): Why did adding Soft Dice Loss ($\lambda=1.0$) boost accuracy from 96.5% to 99.53%?

Answer (Piyush): Mutual Information aligns global tissue intensity distributions but can leave minor boundary errors. Soft Dice loss directly penalizes mask misalignments, forcing the network to align brain tissue edges with sub-voxel precision.

Q4 (Sreehari): How did you achieve a 3,600x inference speedup over classical B-Spline?

Answer (Sreehari): B-Spline runs 1,000 iterative CPU loops per patient (~3 minutes). VoxelMorph shifts all heavy learning offline to training. During inference, it predicts the 3D deformation field in a single forward pass on GPU CUDA cores in 50 milliseconds.

Q5 (Shreeshath): Why did you reorient scans to RAS+ coordinate system?

Answer (Shreeshath): Medical scanners use different axis orientation conventions (e.g. head-first vs feet-first). SimpleITK RAS+ reorientation guarantees all brain volumes face the exact same direction along physical axes.

Q6 (Mahi): What is the main difference between Rigid, Affine, and B-Spline registration?

Answer (Mahi): Rigid allows 6 DOF (translation/rotation). Affine allows 12 DOF (adds scaling/shearing). B-Spline adds a 3D control point grid for local non-linear soft tissue deformation.

Q7 (Piyush): What is the purpose of VecInt diffeomorphic integration?

Answer (Piyush): VecInt uses 7 scaling and squaring steps to integrate stationary velocity fields into displacement fields. This guarantees the warp is smooth, invertible, and physically realistic.

Q8 (Sreehari): Why does the difference map show a bright yellow border around the skull?

Answer (Sreehari): That is an expected physical intensity gap, not a registration error. MRI cannot detect bone (appears dark), while CT shows bone as bright white (+1000 HU). The internal brain tissue shows dark red (near-zero error), confirming true anatomical alignment.

Q9 (Piyush): Why is 3D Elastic Data Augmentation important?

Answer (Piyush): It applies random 3D deformations during training so the network learns to handle arbitrary head posture shifts without overfitting.

Q10 (Sreehari): What are the primary technical limitations of your model?

Answer (Sreehari): The model is trained exclusively on brain scans (cannot register pelvis without retraining), requires rigid pre-alignment for head tilts >30 degrees, and requires prospective radiologist review prior to clinical hospital deployment.

SECTION 7: LIMITATIONS & FUTURE WORK

7.1 Technical Limitations

1. **Single Anatomy (Brain Only):** Trained exclusively on brain MRI-CT pairs; pelvis/thorax require retraining.
2. **Pathology Handling:** Non-pathological dataset; severe tumors or resection cavities require fine-tuning.
3. **Extreme Head Rotations:** Head tilts >30 degrees require rigid pre-alignment prior to VoxelMorph inference.
4. **Clinical Validation:** Dice measures overlap; prospective radiologist review is required before hospital deployment.

7.2 Future Roadmap

1. **Interactive Web Dashboard:** FastAPI + Next.js web application with 3D crosshair slice viewer.
2. **Whole-Body Expansion:** Retrain framework on SynthRAD 2023 Pelvis and Thorax datasets.
3. **DICOM PACS Integration:** Direct hospital PACS integration for automated radiation therapy workflows.